

# CS 432: Databases

## Assignment 8: Query Optimization

### The Eight

Q1) Optimize the query (having Like statement) for searching in the text-type column with the clause -

‘UNION’. (Hint: Query can be stated as searching first\_name’s starting with ‘A’ OR last\_name’s starting with ‘B’)

```

7 #A8 Q1
8 #Insert indexing to improve performance
9 • explain Select username from accounts where account_id in
10 (select account_id from user where city like "Sadar Bazar" union select account_id from vendor where city like "Sadar Bazar" );
11

```

#	id	select_type	table	partition:	type	possible_key:	key	key_len	ref	rows	filtered	Extra
	1	PRIMARY	accounts		index		username_UNIQUE	402		80	100.00	Using where; Using index
	2	DEPENDENT SUBQUERY	user		eq_ref	PRIMARY	PRIMARY	4	func	1	11.11	Using where
	3	DEPENDENT UNION	vendor		eq_ref	PRIMARY	PRIMARY	4	func	1	11.11	Using where
		UNION RESULT	<union2,3>		ALL							Using temporary

```

7 #A8 Q1
8 #Insert indexing to improve performance
9 • Select username from accounts where account_id in
10 (select account_id from user where city like "Sadar Bazar" union select account_id from vendor where city like "Sadar Bazar" );
11

```

Query_ID	Duration	Query
1	0.00009350	SHOW WARNINGS
2	0.00131525	Select username from accounts where account_id in (select account_id from user where city like "Sadar Bazar" union select account_id from vendor where city like "Sadar Bazar" ) LIMIT 0, 1000

After optimization

```

7 #A8 Q1
8 #Insert indexing to improve performance
9 • explain Select username from accounts where account_id in
10 (select account_id from user where match(city) against("Sadar Bazar") union select account_id from vendor where match(city) against( "Sadar Bazar" ) );
11

```

#	id	select_type	table	partition:	type	possible_keys	key	key_len	ref	rows	filtered	Extra
	1	PRIMARY	accounts		index		username_UNIQUE	402		80	100.00	Using where; Using index
	2	DEPENDENT SUBQUERY	user		eq_ref	PRIMARY,account_id_idx1	PRIMARY	4	func	1	11.11	Using where
	3	DEPENDENT UNION	vendor		eq_ref	PRIMARY,account_id_idx3	PRIMARY	4	func	1	11.11	Using where
		UNION RESULT	<union2,3>		ALL							Using temporary

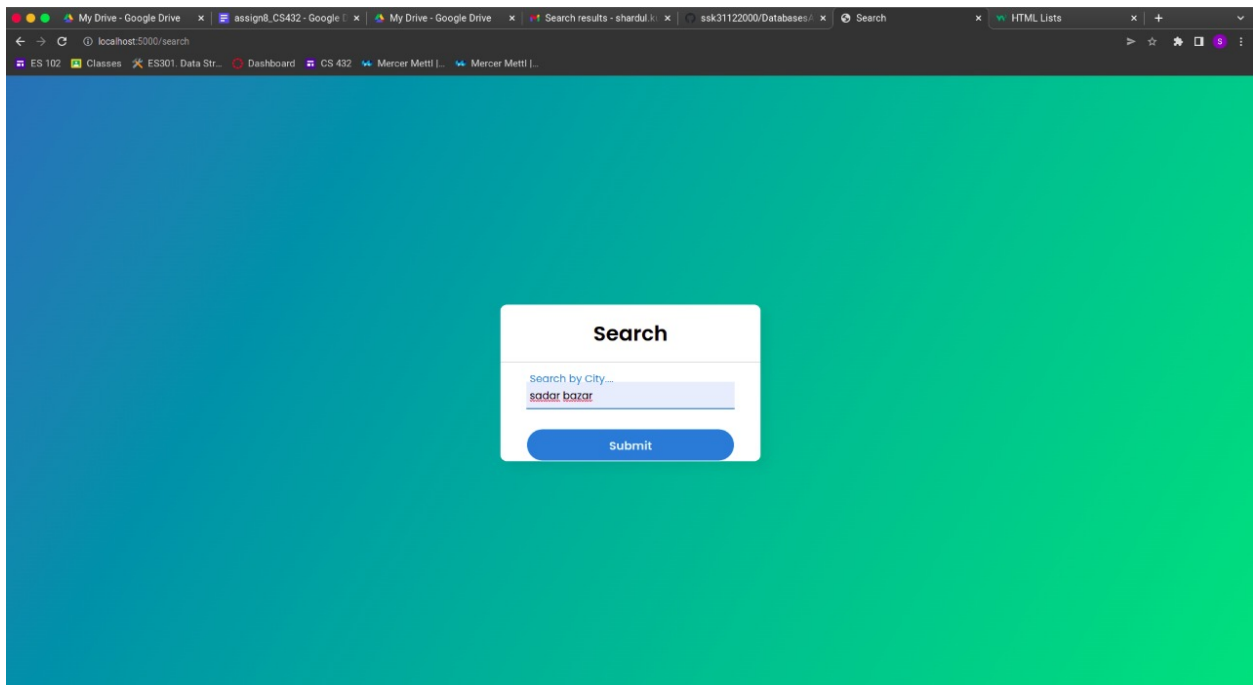
```

7 #A8 Q1
8 #Insert indexing to improve performance
9 • Select username from accounts where account_id in
10 (select account_id from user where match(city) against("Sadar Bazar") union select account_id from vendor where match(city) against( "Sadar Bazar" ) );
11

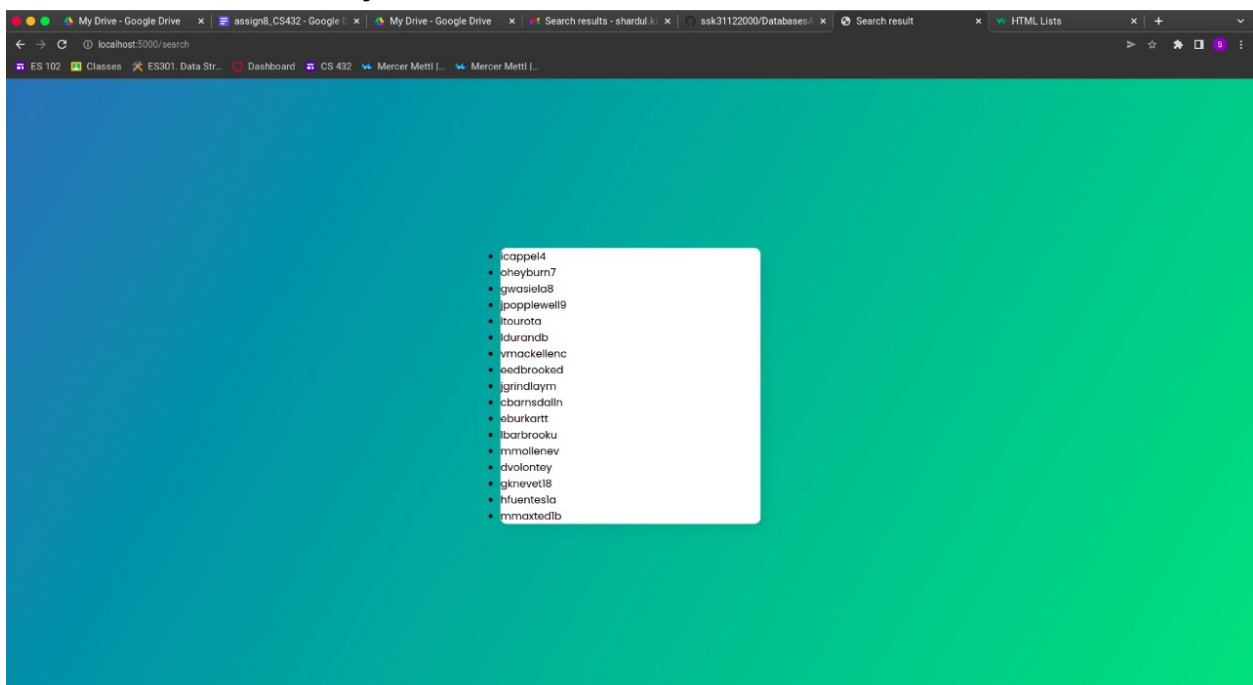
```

Query_ID	Duration	Query
9	0.00083425	explain Select username from accounts where account_id in (select account_id from user where match(city) against("Sadar Bazar") union select account_id from vendor where match(city) against(...
10	0.00203875	Select username from accounts where account_id in (select account_id from user where match(city) against("Sadar Bazar") union select account_id from vendor where match(city) against(... LIMIT 0, 1000

We have applied full text indexing on the union of subqueries of searching city names .But we got more execution time after indexing compared to the original query.  
Search by the locality/city.



Users/Username in the city as result of before search.



(Q:2)Report the number of scans for optimizing the search through a column having text values with a chosen text pattern (Any text pattern of your choice). (5M)

❖ Finding the name of the city ending with “Bazar”

Original query:

```
select * from vendor WHERE city like '%Bazar%';
```

Optimised query:

```
ALTER TABLE vendor  
ADD FULLTEXT(city);
```

```
select * from vendor WHERE MATCH(vendor.city) AGAINST ('Bazar')
```

Stats:

```
1 • create database ecom;  
2 • use ecom ;  
3 • select * from vendor;  
4 • explain select * from vendor WHERE city like '%Bazar%';  
5
```

Result Grid

Filter Rows:

Export:

Wrap Cell Content: [FA](#)

	id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
▶	1	SIMPLE	vendor	NULL	ALL	NULL	NULL	NULL	NULL	20	11.11	Using where

	account_id	warehouse_details	pincode	city	country	phone_number
▶	53	0453 Carioca Parkway	281002	Sadar Bazar	India	5483257278
	57	5 Knutson Way	281002	Sadar Bazar	India	1584164283
	59	523 Porter Circle	281002	Sadar Bazar	India	8659122018
	67	74 Corscot Way	281002	Sadar Bazar	India	8425774684

Number of scanned rows=20

Execution

time=

0.0004510

0

### Optimised Query:

```
1 • create database ecom;  
2 • use ecom ;  
3 • ALTER TABLE vendor  
4   ADD FULLTEXT(city);  
5 • explain select * from vendor WHERE MATCH(vendor.city) AGAINST ('Bazar')
```

Result Grid   Filter Rows:   Export:   Wrap Cell Content:												
	id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	1	SIMPLE	vendor	NULL	fulltext	city	city	0	const	1	100.00	Using where; Ft_hints: sort

Execution time=

0.0001420

Number of scanned rows=1

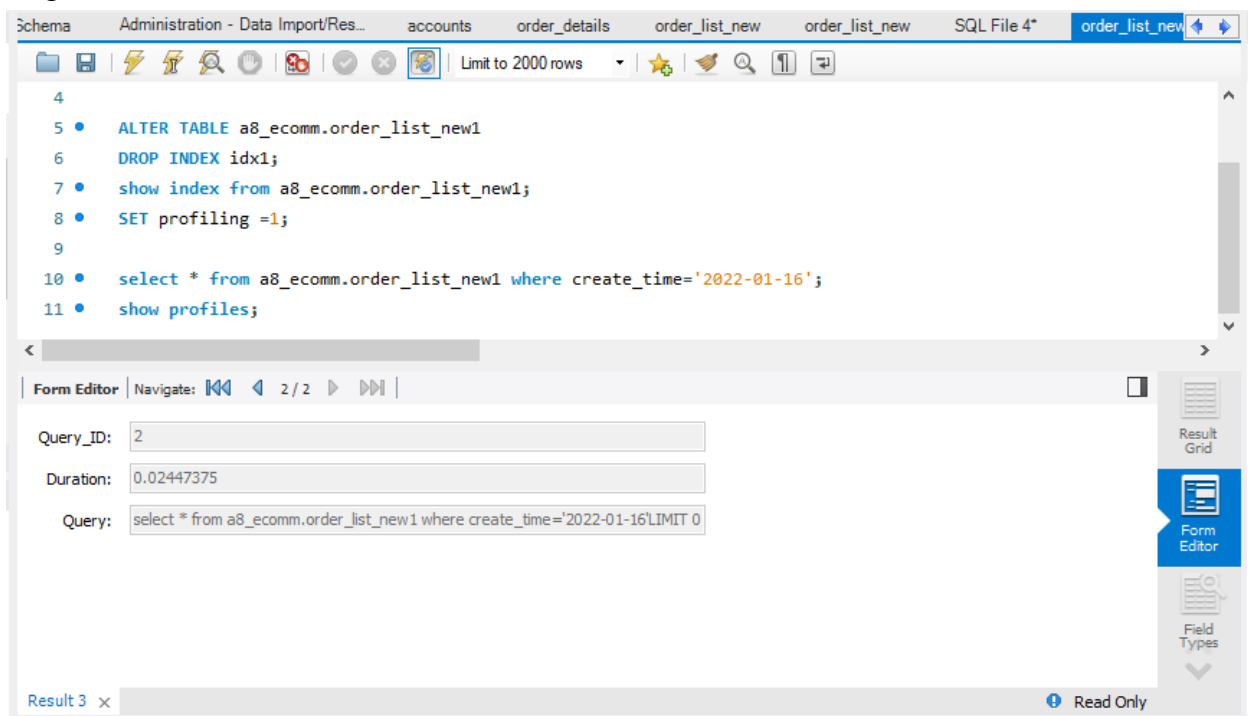
I have used full text search indexing to find the city column names ending with “bazar” to optimise the query .

**Q3. Suppose there are limited data entries (Say 50 users' details), ideally it would be better if we change the id's data type to TINYINT (the signed range of TINYINT is from -128 to 127). Similarly, change the data type for either one from [phone\_number, pin\_code, address, or any other specific to your statement].**

```
12 #A8 Q3
13 • alter table order_list modify order_quantity tinyint; #Since quantity will be less than 127 (assumed)
```

**Q4.**

Original:



Some are below snapshots of trials with execution time for the various combinations with indexing type, groupby, order by etc operations.

Order By Create\_time Ascending

```

8 • select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY create_time;
9 • show profiles;

```

Form Editor | Navigate: ⏮ ⏪ 7 / 7 ⏩ ⏭

Query\_ID: 7

Duration: 0.00302125

Query: select \* from a8\_ecomm.order\_list\_new1 where create\_time='2022-01-16' ORDER BY create\_time;

Result Grid

Form Editor

Field Types

### Order By Create\_time Descending

```

8 • select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY create_time DESC;
9 • show profiles;

```

Form Editor | Navigate: ⏮ ⏪ 8 / 8 ⏩ ⏭

Query\_ID: 8

Duration: 0.00714750

Query: select \* from a8\_ecomm.order\_list\_new1 where create\_time='2022-01-16' ORDER BY create\_time DESC;

Result Grid

Form Editor

Field Types

### ORDER BY YEARcreate\_time;

```

6 -- DROP INDEX idx1;
7 • show index from a8_ecomm.order_list_new1;
8 • select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY YEARcreate_time;
9 • show profiles;

```

Form Editor | Navigate: ⏮ ⏪ 10 / 10 ⏩ ⏭

Query\_ID: 10

Duration: 0.00369550

Query: select \* from a8\_ecomm.order\_list\_new1 where create\_time='2022-01-16' ORDER BY YEARcreate\_time;

Result Grid

Form Editor

Field Types

### ORDER BY MONTHcreate\_time;

```

8 • select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY MONTHcreate_time;
9 • show profiles;

```

Form Editor | Navigate: ⏮ ⏪ 11 / 11 ⏩ ⏭

Query\_ID: 11

Duration: 0.00328475

Query: select \* from a8\_ecomm.order\_list\_new1 where create\_time='2022-01-16' ORDER BY MONTHcreate\_time;

Result Grid

Form Editor

ORDER BY MONTHcreate\_time, YEARcreate\_time ;

```

8 • select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY MONTHcreate_time, YEARcreate_time;
9 • show profiles;

```

Form Editor | Navigate: ⏮ ⏪ 14 / 14 ⏩ ⏭

Query\_ID: 14

Duration: 0.00271600

Query: select \* from a8\_ecomm.order\_list\_new1 where create\_time='2022-01-16' ORDER BY MONTHcreate\_time, YEARcreate\_time;

Result Grid

Form Editor

Unique index of order\_id and orderby MONTHcreate\_time, YEARcreate\_time

```

3 • create unique index order_idx on a8_ecomm.order_list_new1(order_id);
4
5 • ALTER TABLE a8_ecomm.order_list_new1
6 DROP INDEX order_idx;
7 • show index from a8_ecomm.order_list_new1;
8 • select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY MONTHcreate_time, YEARcreate_time;
9 • show profiles;

```

Form Editor | Navigate: ⏮ ⏪ 15 / 15 ⏩ ⏭

Query\_ID: 22

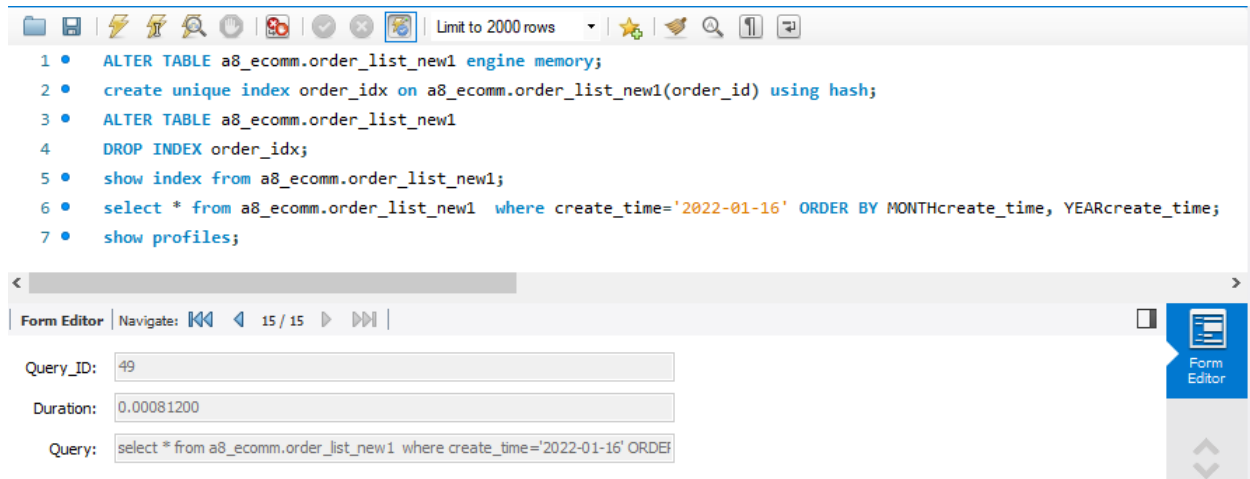
Duration: 0.00166050

Query: select \* from a8\_ecomm.order\_list\_new1 where create\_time='2022-01-16' ORDER BY MONTHcreate\_time, YEARcreate\_time;

Result Grid

Form Editor

order list new1 21 Result 22 x Read Only



Concluding from above all the snapshots we do the following steps for optimizing query as follows:

- 1) First we alter the table by using engine memory by adding the following query:

```
ALTER TABLE a8_ecomm.order_list_new1 engine memory;
```

- 2) Then we create the unique index named “order\_idx” using the column order\_id by indexing method of btree.

```
create unique index order_idx on a8_ecomm.order_list_new1(order_id) using btree;
```

- 3) Our main query is finding certain date, for optimizing we split the column of create\_time (date) into dd (day) | mm(month)| y(year) seperated by columns. Year column of create\_time is named as “YEARcreate\_time” and Month column of create\_time is named as “MONTHcreate\_time”

```
select * from a8_ecomm.order_list_new1 where create_time='2022-01-16' ORDER BY MONTHcreate_time, YEARcreate_time;
```

Further we order by YEARcreate\_time, for search of create\_time = 2022-01-16 and obtain the following results with minimum duration of execution.

We have iterated various options as follows for optimization.

Indexing Type + Indexing Method	GROUP BY	ORDER BY	DURATION
Btree + Unique	-	MONTHcreate_time	~0.00067125
Btree + Unique	-	YEARcreate_time	~0.00133925
Btree + Unique	-	MONTHcreate_time,	~0.00100400



		YEARcreate_time	
Btree + Unique	-	create_time	~0.005
Hashing + Unique	-	MONTHcreate_time	~0.003
Hashing + Unique	-	YEARcreate_time	~0.006
Hashing + Unique	-	MONTHcreate_time, YEARcreate_time	~0.00081200

In the above iterations we cannot use fulltext since order\_id is integer which is the primary key.

Q5.

Original Table				Added Intentional Null Values.			
	category_id	category	description		category_id	category	description
▶	5001	Electronics	Endosc destru bile les	▶	5001	Electronics	Endosc destru bile les
	5002	Smartphones	Occlude leg vein NEC		5002	Smartphones	Occlude leg vein NEC
	5003	Fashion	Ins/rep 1 pul gen,rechrg		5003	Fashion	NULL
	5004	Beauty	Limb shorten-metacar/car		5004	Beauty	Limb shorten-metacar/car
	5005	Books	Sm bowel endoscopy NEC		5005	Books	Sm bowel endoscopy NEC
	5006	Toys	Pericardiocentesis		5006	Toys	Pericardiocentesis
	5007	Home Decor	Bronchial operation NEC		5007	Home Decor	NULL
	5008	Kitchen	Dilat frontonasal duct		5008	Kitchen	NULL
	5009	Laptops	Removal brain stim lead		5009	Laptops	Removal brain stim lead
	5010	Shoes	Suture anal laceration		5010	Shoes	Suture anal laceration

Query Before:

```
SELECT count(*) FROM a8_ecomm.cat_null;
```

Answer:

	count(*)
▶	10

Query Added:

```
SELECT count(*) FROM a8_ecomm.cat_null where description is NOT NULL;
```

Answer without including NULL:

	count(*)
▶	7

## Q6.

```
C:\Users\tanis>mysql -u root -p
```

```
mysql> SHOW STATUS LIKE "qcache%";
Empty set (0.14 sec)
```

```
mysql> SHOW VARIABLES LIKE 'have_query_cache';
+-----+-----+
| Variable_name | Value |
+-----+-----+
| have_query_cache | NO    |
+-----+-----+
1 row in set (0.11 sec)
```

As can be seen in the images above the query cache doesn't work. The query cache is deprecated as of MySQL 5.7.20, and is removed in MySQL 8.0.

Tried overwriting the MySQL config file, but failed.

✗ Failed to save 'my.ini': EPERM: operation not permitted, open  
'C:\Users\tanis\AppData\Local\Temp\c4b95a1ac802053ca9335d0c0  
9660ac7\status'

The query cache has been disabled because it does not scale with high-throughput workloads on multi-core machines. For more information on this we can refer to a blog by Morgan tocker, [click here](#) for link.

#### Q7.

##### Queries used:

- select username, password from accounts left join admin on (accounts.account\_id = admin.account\_id) where admin.authorisation = 1;
- select username, password from accounts where account\_id = (select account\_id from admin where authorisation = 1 and admin.account\_id = accounts.account\_id);

✓	320	16:37:54	select * from accounts left join admin on (accounts.account_id = admin.account_id) where admin.authorisation = 1 LIMIT 0, 1000	4 row(s) returned	0.015 sec / 0.000 sec
✓	321	16:42:15	select username, password from accounts where account_id = (select account_id from admin where authorisation = 1 and admin.account_id = accounts.account_id) LIMIT 0, 1000	4 row(s) returned	0.109 sec / 0.000 sec

From the above operation, we see join operation consumes less time over subqueries.

Disadvantages of using multiple join operation:

1. Multiple Joins requires the server to do more work, thus consuming more time to retrieve it.
2. Different types of joins can create confusion, as each join gives out different results.
3. Usage of incorrect joins can result in performance degradation and inaccurate results.
4. Joins are not easy to read as subqueries.

#### Q.8

Unoptimized query:

```

7 • explain SELECT city FROM
8 (SELECT user_id, total_price,city,create_time FROM user RIGHT JOIN order_details ON user.account_id=order_details.user_id
9 WHERE total_price=(SELECT MAX(total_price) FROM order_details))
10 AS temp
11 WHERE create_time=
12 (SELECT MAX(create_time) FROM user RIGHT JOIN order_details ON user.account_id=order_details.user_id WHERE total_price=(SELECT MAX(total_price) FROM order_details));
13
14 • SET profiling=1;
15 • SHOW profiles;

```

result Grid
Filter Rows:
Export:
Wrap Cell Content:

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	order_details		ALL					9	11.11	Using where
1	PRIMARY	user		eq_ref	PRIMARY,user_idx	PRIMARY	4	a8_ecom.order_details.user_id	1	100.00	
4	SUBQUERY	order_details		ALL					9	11.11	Using where
4	SUBQUERY	user		eq_ref	PRIMARY,user_idx	PRIMARY	4	a8_ecom.order_details.user_id	1	100.00	Using index
5	SUBQUERY	order_details		ALL					9	100.00	
3	SUBQUERY	order_details		ALL					9	100.00	

Query\_ID:

16

Duration:

0.00219850

Query:

SELECT city FROM (SELECT user\_id, total\_price,city,create\_time FROM user RIGH

Number of scans=9+1+9+1+9+9=38  
 Execution time=0.00219850 sec

Optimized Query:

```

1 • explain SELECT city FROM
2   (SELECT user_id, total_price,city,create_time FROM user RIGHT JOIN order_details ON user.account_id=order_details.user_id
3    WHERE total_price=(SELECT MAX(total_price) FROM order_details))
4   AS temp
5   ORDER BY create_time DESC LIMIT 1;
6
7 • show profiles;

```

result Grid

Filter Rows:

Export:

Wrap Cell Content:

IA

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	order_details	NULL	ALL	NULL	NULL	NULL	NULL	9	11.11	Using where; Using filesort
1	PRIMARY	user	NULL	eq_ref	PRIMARY,user_idx	PRIMARY	4	a8_ecom.order_details.user_id	1	100.00	NULL
3	SUBQUERY	order_details	NULL	ALL	NULL	NULL	NULL	NULL	9	100.00	NULL

Query\_ID: 17

Duration: 0.00043400

Query: SELECT city FROM (SELECT user\_id, total\_price,city,create\_time FROM user RIGHT JOIN order\_details ON user.account\_id=order\_details.user\_id WHERE total\_price=(SELECT MAX(total\_price) FROM order\_details)) AS temp ORDER BY create\_time DESC LIMIT 1;

Number of scans=9+1+9=19  
 Execution time=0.00043400 sec